

CS 354 - Machine Organization & Programming

Tuesday Feb 14 and Thursday Feb 16, 2023

Midterm Exam - Thursday, February 23th, 7:30 - 9:30 pm

- ♦ Room: Students will be assigned a room based and sent email with that room
- ♦ UW ID required
- ♦ #2 pencils required
- ♦ closed book, no notes, no electronic devices (e.g., calculators, phones, watches)
- ♦ see “Midterm Exam 1” on course site Assignments for topics
- ♦ see sample exam cover page (last page of this outline)

Project p2A: Due on or before Friday, Feb 17th

Project p2B: Due on or before Friday, Feb 24th

Homework hw1: Due on Monday Feb 13th (solution available Wed morning)

Homework hw2: Due on Monday Feb 20th (solution available Wed morning)

Last Week

Command-line Arguments Recall 2D Arrays 2D Arrays on the Heap 2D Arrays on the Stack 2D Arrays: Stack vs. Heap	Array Caveats Meet Structures Nesting in Structs and Arrays of Structs Passing Structures
--	--

This Week

Pointers to Structures (from L6) Standard & String I/O in <code>stdio.h</code> File I/O in <code>stdio.h</code> Copying Text Files Three Faces of Memory	Virtual Address Space C's Abstract Memory Model Meet Globals and Static Locals Where Do I Live? Linux: Processes and Address Spaces Exam Sample Cover Page
Next Week: The Heap & Dynamic Memory Allocators (p3) Read: B&O 9.1, 9.2, 9.9.1-9.9.6 9.1 Physical and Virtual Addressing 9.2 Address Spaces 9.9 Dynamic Memory Allocation 9.9.1-9.9.6	

Standard and String I/O in `stdio.h`

Standard I/O

Standard Input

`getchar` //reads 1 char

don't use for string `gets` //reads 1 string ending with a newline char, BUFFER MIGHT OVERFLOW

`int scanf(const char *format_string, &v1, &v2, ...)`

reads formatted input from the console keyboard

returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs

format string contains format chars and chars to skip

format specifiers `%d, %f, %c, %i, %p, %s`

←
octal, decimal, hex

whitespace input separator (space, tab, newline)

leading whitespace is skipped

Standard Output

`putchar` //writes 1 char

`puts` //writes 1 string

`int printf(const char *format_string, v1, v2, ...)`

writes formatted output to the console terminal window

returns number of characters written, or a negative if error

format string format specifiers and chars to print

use `\n` to flush output buffer

Standard Error

`void perror(const char *str)`

writes formatted error output to the console terminal window

String I/O

input
`int sscanf(const char *str, const char *format_string, &v1, &v2, ...)`

reads formatted input from the specified `str`

returns number of characters read, or a negative if error

`int sprintf(char *str, const char *format_string, v1, v2, ...)`

writes formatted output to the specified `str`

returns number of characters written, or a negative if error

File I/O in `stdio.h`

Standard I/O Redirection IN LINUX

//file extensions (.txt, .exe) rarely used in linux

`a.out < input_file` `> output_file` //overwrites `output_file`
`>> output_file` //appends `output_file`

File I/O in C p2A, p2B

File Input

`fgetc/getc`, `ungetc` //reads 1 char at a time from a file
`fgets` //reads 1 string terminate with a newline char or EOF
`int fscanf(FILE *stream, const char *format_string, &v1, &v2, ...)`
reads formatted input from the specified `stream`
returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs

File Output

`fputc/putc` //writes 1 char at a time
`fputs` //writes 1 string
`int fprintf(FILE *stream, const char *format_string, v1, v2, ...)`
writes formatted output to the specified `stream`
returns number of characters written, or a negative if error

Predefined File Pointers

`stdin` is console keyboard
`stdout` is console terminal window
`stderr` is console terminal window, second stream for errors

`printf("hello\n");` \equiv `fprintf(stdout, "hello\n");`

Opening and Closing

"r" read
"w" write

`FILE *fopen(const char *filename, const char *mode)`
opens the specified `filename` in the specified `mode`
returns file pointer to the opened file's descriptor, or NULL if there's an access problem
`int fclose(FILE *stream)`
flushes the output buffer and then closes the specified `stream`
returns 0, or EOF if error

Copying Text Files

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {

    if (argc != 3) {
        fprintf(stderr, "Usage: copy inputfile outputfile\n");
        exit(1);
    }

    FILE *ifp = fopen(argv[1], "r");
    if (ifp == NULL) {
        fprintf(stderr, "Can't open input file %s!\n", argv[1]);
        exit(1);
    }

    FILE *ofp = fopen(argv[2], "w");
    if (ofp == NULL) {
        fprintf(stderr, "Can't open output file %s!\n", argv[2]);
        fclose(ifp);
        exit(1);
    }

    const int bufsize = 257; //WARNING: assumes lines <= 256 chars
    char buffer[bufsize]; //STACK
    while (fgets(buffer, bufsize, ifp) != NULL) {
        fputs(buffer, ofp);
    }
    fclose(ifp);
    fclose(ofp);

    return 0;
}
```

Three Faces of Memory


* **Abstraction:** manage complexity by focusing on relevant details

Process View = Virtual Memory

Goal: provide a simple view

virtual address space (VAS): an illusion by the OS that each process has its own contiguous memory space

virtual address: simulated address that a process generates



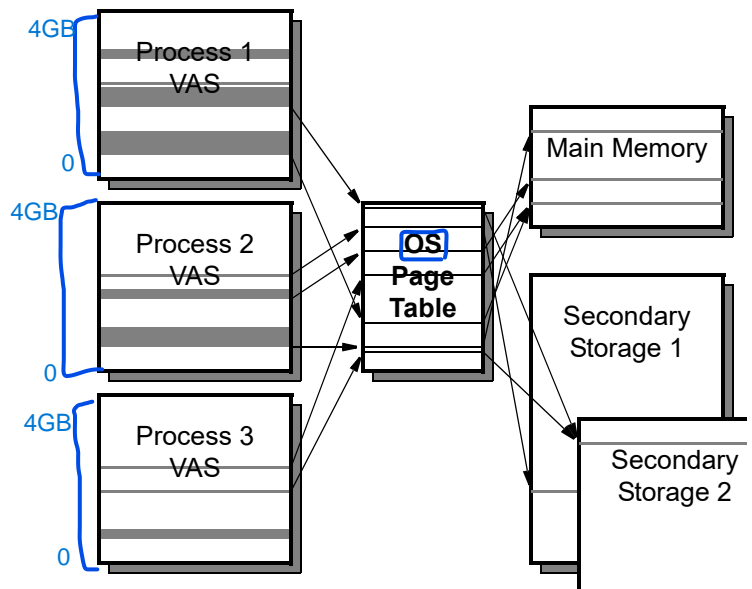
System View = Illusionist (CS 537)

Goal: make memory shareable & secure

pages: 4K (kilobytes)
mem block managed by OS

page table: OS data struct to track page allocation

VAS -> PAS for each process
(virtual to physical address)

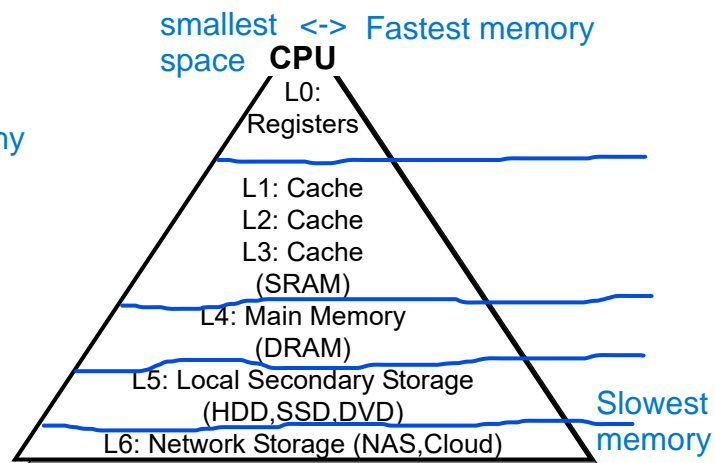


Hardware View = Physical Memory

Goal: Keep CPU busy

physical address space (PAS): Multi-level hierarchy ensures frequently accessed data is close to CPU

physical address: actual address used to access machine's memory



Virtual Address Space (IA-32/Linux)

32-bit Processor = 32-bit Addresses => $2^{32} = 4,294,967,296 = 4\text{GB}$ Address Space

11111111111111111111111111111111 = 0xFFFFFFFF

address space: a range of valid addresses for a process

process: a running process

11000000000000000000000000000000 = 0xC0000000

kernel: memory resident part of OS

user process: process that is not in KERNEL. "our" processes


✳ Every user process has this simple view of memory

for DLL's, file I/O, large heap allocation

*****REMEMBER CODE, DATA, HEAP, AND STACK

00001000000001001000000000000000 = 0x08048000

00000000000000000000000000000000 = 0x00000000



C's Abstract Memory Model

1. **CODE** Segment

Contains: the program

.text section instructions

ro =
read only .rodata section string literals

Lifetime: entire program's execution

Initialization: from EOF by LOADER

Access: READ-ONLY

2. **DATA** Segment

Contains: global variables and static local

Lifetime: entire program's execution

Initialization: from EOF by LOADER

.data section from code, for variables initialized to non-zero values

.bss section from code, for variables that are initialized to zero or uninitialized
Block Started Symbol

Access: read/write

3. **HEAP** (AKA Free Store)

Contains: memory allocated and freed by the programmer during execution

Lifetime: managed by the programmer, malloc/calloc/realloc until free

Initialization: none by default

Access: read/write

4. **STACK** (AKA Auto Store)

Contains: memory in stack frames, auto allocated and freed by function call/return

stack frame (AKA activation record) S.F., A.R., parameters, local variables, temp variables, more...

Lifetime: from declaration to end of scope, often a "closing brace" }

Initialization: none by default

Access: read/write

Meet Globals and Static Locals

What?

A global variable is

- ◆ declared outside of a function
- ◆ accessible to all functions in a source file
- ◆ allocated in **DATA** segment

A static local variable is

- ◆ declared in a function with static modifier
- ◆ accessible only within the function
- ◆ allocated in **DATA** segment

Why?

for storage that exists during entire program

✱ *In general, global variables should not be used!*

Instead use local variables that are passed to the callee function

How?

```
#include <stdio.h>
int g = 11; //global

void f1(int p) { //parameter (very similar to local variable)
    static int x = 22; //static local
    x = x + p * g;
    printf("%d\n", x);
}

int main(void) {
    f1(g);
    g = 2;
    int g = 1; //local variable that shadows global variable
    f1(g);
    return 0;
}
```

shadowing: when a local variable blocks access to a global variable of same name

✱ *Avoid shadowing; don't use the same identifier for local and global variables*

Where do I live? which segment of VAS contains memory for each variable

→ Identify the segment (and section) for each memory allocation in the code below.

```
#include <stdio.h>
#include <stdlib.h>

int gus = 14;    DATA .data
int guy;        DATA .bss

int madison(int pam) {    STACK

    static int max = 0; DATA .bss
    int meg[] = {22,44,88};    STACK
    int *mel = &pam;
    max = gus--;
    return max + meg[1] + *mel;
}

int *austin(int *pat) {    STACK

    static int amy = 33; DATA .data
    int *ari = malloc(sizeof(int)*44);    STACK    HEAP
    gus--;
    *ari = *pat;
    return ari;
}

int main(int argc, char *argv[]) {    STACK

    int vic[] = {33,66,99};    SAA
    int *wes = malloc(sizeof(int));    STACK    HEAP
    *wes = 55;
    guy = 66;
    free(wes);
    wes = vic;
    wes[1] = madison(guy);
    wes = austin(&gus);
    free(wes);
    printf("Where do I live?");
    return 0;
}
```

* *Arrays, structs, and variables* can live in DATA, HEAP, STACK

Pointer var can store any address, SEG FAULT if dereference outside of your memory

Linux: Processes and Address Spaces

Process and Job Control

- ♦ Linux is a multitasking operating system where you can run multiple processes concurrently

ps snapshot of user processes

jobs only user process

& put process in background

ctrl+z suspend running process

bg put a suspended process in the background

fg bring suspended process to foreground

ctrl+c stop a running process

Program Size

size <executable or object_file> a.out, prog.

```
$gcc -m32 myProg.c
```

```
$size a.out
```

text	data	bss	dec	hex	filename
1029	276	4	1309	51d	a.out
CODE	DATA		TOTAL		

Virtual Address Space Maps

- ♦ Linux enables you to see the virtual address space of each process

```
$pmap <pid_of_process>
```

```
$cat /proc/<pid_of_process>/maps
```

```
$cat /proc/self/maps
```

cat /proc:

[] Lec 001 9:30am TR		
[] Lec 002 1:00pm TR		
Lecture	Print Netid	Print Name (first last)

Computer Sciences 354
Midterm Exam 1 Secondary
Thursday, October 6th, 2022
60 points (15% of final grade)
Instructors: Debra Deppeler

1. MARK an X in box by your lecture number above.
2. PRINT your NET ID (UW login name not your photo id number) in box above.
3. PRINT your first and last name in box above.
4. FILL-IN all fields and their bubbles on the scantron form (use # 2 pencil).
 - (a) LAST NAME -fill in your last (family) name starting at leftmost column.
 - (b) FIRST NAME -fill in first five letters of your first (given) name.
 - (c) IDENTIFICATION NUMBER is your UW Student ID number.
 - (d) Under ABC of SPECIAL CODES, write your lecture number as a three digit value 001 or 002.
 - (e) Under F of SPECIAL CODES, write the number 2 for Secondary and fill in the number (2) bubble.
5. DOUBLE-CHECK THAT YOU HAVE FILLED IN ALL ID FIELDS and that you have FILLED IN ALL CORRESPONDING BUBBLES ON SCANTRON.
6. Taking this exam indicates that you agree: to not write answers in large letters and to keep your answers covered; to not view or use another's work or any unauthorized devices in any way; to not make any type of copy of any portion of this exam; and that you understand that being caught doing any of these actions, or other actions that permit any student to submit work that is not wholly their own will result in automatic failure of the exam and possible failure of the course. Penalties are reported to the Deans Office for all involved.

Parts	Number of Questions	Question Format	Possible Points
I	10	Simple Choice	20
II	10	Multiple Choice	30
III	2	Written	10
	22	Total	60

Assumptions unless instructions explicitly state otherwise:
addresses and integers are 4 bytes.
code questions are about C and IA-32/x86 assembly code on our Linux platform.

Reference: Powers of 2

$$2^5 = 32, 2^6 = 64, 2^7 = 128, 2^8 = 256, 2^9 = 512, 2^{10} = 1024$$

$$2^{10} = K, 2^{20} = M, 2^{30} = G$$

$$2^A * 2^B = 2^{A+B}, 2^A / 2^B = 2^{A-B}$$

Turn off and put away all electronic devices and
wait for the proctor to signal the start of the exam.