

CS 354 - Machine Organization & Programming

Tuesday, March 7 and Thursday March 9, 2023

Print paper copies of this outline for best use.

Week 07 Activity: Heap Practice Assignment

Project p3: DUE on or before Friday March 25

Homework 3: DUE on or before Friday March 11

Last Week

Placement Policies (finish) Free Block - Too Large/Too Small Coalescing Free Blocks Free Block Footers	Explicit Free list Explicit Free List Improvements --- Heap Caveats Memory Hierarchy
---	--

This Week: MEMORY MANAGEMENT via CACHING blocks of memory for fast access

Heap Caveats Memory Hierarchy Locality & Caching Bad Locality Caching: Basic Idea & Terms Designing a Cache: Blocks Rethinking Addressing	Designing a Cache: Sets and Tags Basic Cache Lines Basic Cache Operation Basic Cache Practice
Next Week after Spring Break: Vary cache set size and Cache Writes B&O 6.4.3 Set Associative Caches 6.4.4 Fully Associative Caches 6.4.5 Issues with Writes 6.4.6 Anatomy of a Real Cache Hierarchy 6.4.7 Performance Impact of Cache Parameters	

DO NOT TAKE THIS WEEK OFF if you plan to take Spring Break off

Note: p4A and p4B will be released on Monday after Spring Break

Get p3 done this week. It is possible and will make for an actual Spring Break for you!

p3 - implement and test alloc (partA) and free (partB) by Monday and submit progress

p3 - implement coalesce by Wednesday and submit progress

p3 - complete testing and debugging by Friday and complete final submission

Heap Caveats

Consecutive heap allocations don't result in contiguous payloads!

→ Why?

- Payloads are interspersed with heap structure and possibly padding.
- Placement policies & heap structure can scatter allocations throughout heap.

Don't assume heap memory is initialized to 0!

OS initially clears heap pages for security,
but your recycled heap memory will have your old data
unless you use calloc()

Do free all heap memory that your program allocates!

→ Why are memory leaks bad?

They slowly kill your program's performance by cluttering heap with garbage blocks.
Bad leaks could ultimately consume your entire heap!

→ Do memory leaks persist when a program ends?

No, heap pages are returned to the OS for other uses.

Don't free heap memory more than once!

→ What is the best way to avoid this mistake? NULL freed pointers.

Don't read/write data in freed heap blocks!

→ What kind of error will result? Intermittent error

Don't change heap memory outside of your payload!

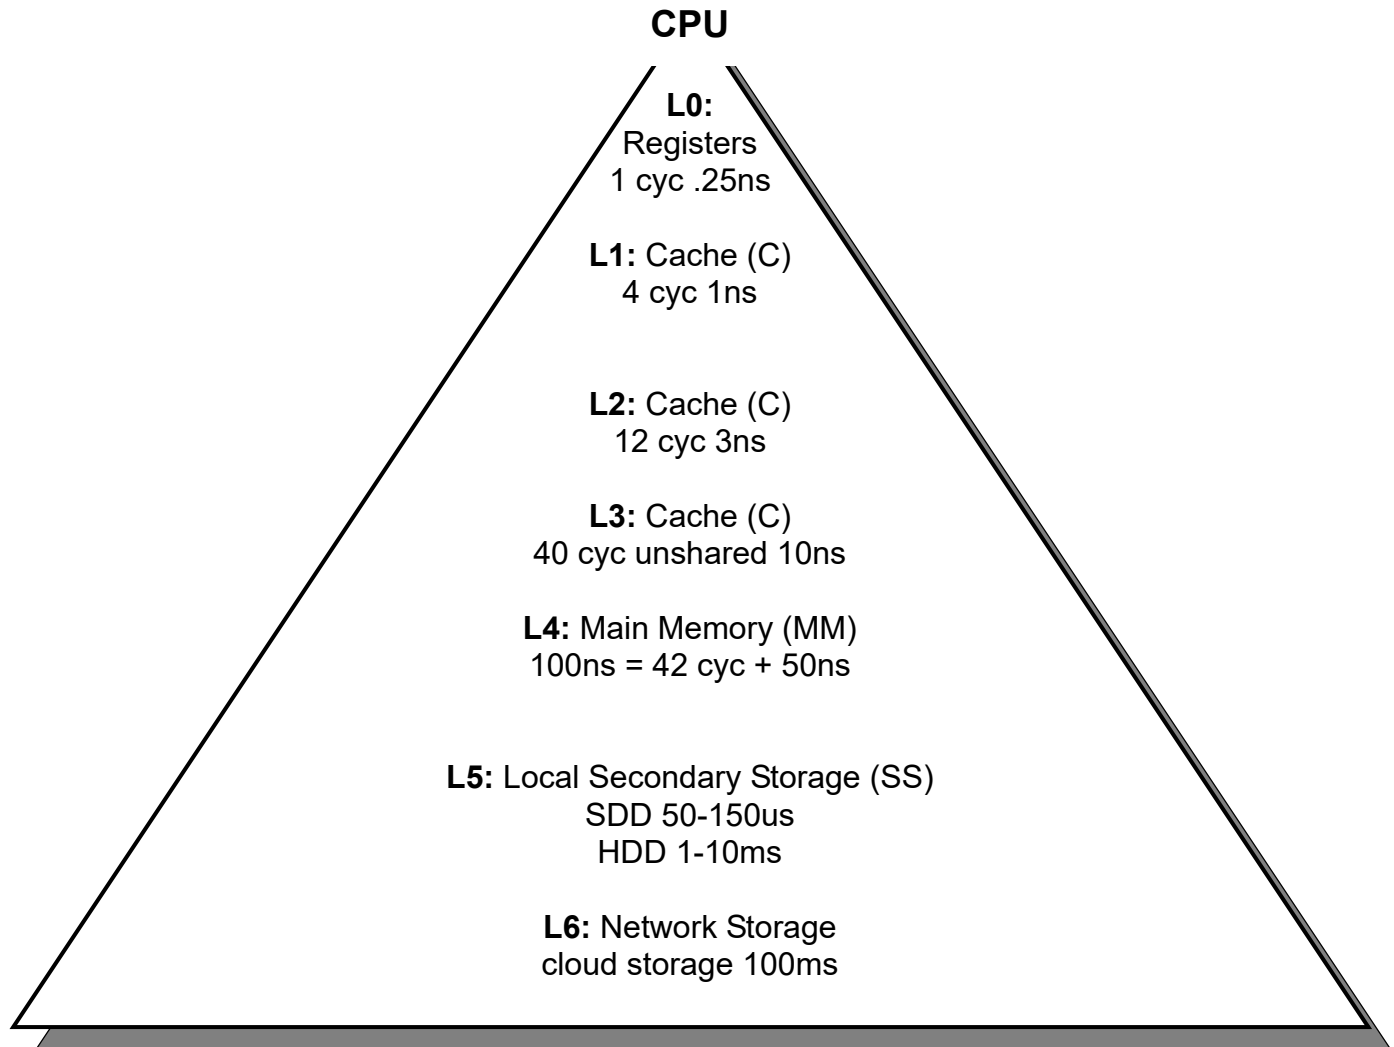
→ Why? You'll trash the heap's internal structure and/or another block's payload.

Do check if your memory intensive program has run out of heap memory!

→ How? Check that allocator's return value is not NULL

Memory Hierarchy

✱ *The memory hierarchy gives the illusion of having lots of fast memory.*



Cache

is a smaller faster mem that acts as a staging area for data stored in a larger slower mem

Memory Units

word: size used by CPU transfer between L1 & CPU

block: size used by C transfer between C levels & MM

page: size used by MM transfer between MM & SS

Memory Transfer Time: [https://simple.wikipedia.org/wiki/Orders_of_magnitude_\(time\)](https://simple.wikipedia.org/wiki/Orders_of_magnitude_(time))

cpu cycles: used to measure time

latency: memory access time (delay)

Locality & Caching

What?

temporal locality: when a recently accessed memory location is repeatedly accessed in the near future

```
int i = 0; i < size; i++
```

spatial locality: when a recently accessed memory location is followed by nearby memory locations being accessed in the near future

locality is designed into Hardware, OS, apps

Example

```
int sumArray(int a[], int size, int step) {  
    int sum = 0;  
    for (int i = 0; i < size; i += step)  
        sum += a[i];  
    return sum;  
}
```

→ List the variables that clearly demonstrate temporal locality. *i, sum, size, step*

→ List the variables that clearly demonstrate spatial locality. *a (if step is small)*

stride: unit in word between sequential accesses

good spatial locality is ~1 word per stride

* The caching system uses locality to predict what the cpu will need in the near future.

↶ both types

How? The caching system anticipates...

temporal: anticipates data will be reused so it copies and saves in cache

spatial: anticipates nearby data will be used so it copies a BLOCK to the cache

cache block: unit of memory transferred between main memory and cache levels

* Programs with good locality run faster since they work better with the caching system!

Why? Programs with good locality maximize use of memory at top of memory hierarchy

Bad Locality

Why is this code bad?


```
int a[ROWS][COLS];
```

```
for (int c = 0; c < COLS; c++)
  for (int r = 0; r < ROWS; r++)
    a[r][c] = r * c;
```

→ How would you improve the code to reduce stride?
rows first then columns, flip for loops


Key Questions for Determining Spatial Locality:

1. What does the memory layout look like for the data?
2. What is the stride of the code across the data? number of columns, smaller stride is better




Why is this code bad?

```
struct {
  float rgb[3];
  float hsl[3];
} image[HEIGHT][WIDTH];
```



```
for (int v = 0; v < 3; v++) v = 0, 1
for (int c = 0; c < WIDTH; c++) c = 0, 1, 2, 0
for (int r = 0; r < HEIGHT; r++) { r = 0, 1, 0, 1, 0, 1, 0
  image[r][c].rgb[v] = 0;
  image[r][c].hsl[v] = 0;
}
```

→ How would you improve the code to reduce stride?



Good or bad locality?

- ♦ Instruction Flow:
sequencing? good spatial locality

(if/switch stmts) selection? bad s and t locality

(loops) repetition? good spatial if data is contiguous and stride is small,
good temporal for loop counters, sentinel values, step index

- ♦ Searching Algorithms:

linear search array good spatial, temporal for match value

linked list bad S

binary search array bad spatial

Caching: Basic Idea & Terms

Assume: Memory is divided into 32 byte blocks and all blocks are already in main memory.
Cache L1 has 4 locations to store blocks and L2 has 16 locations to store blocks.

→ Update the memory hierarchy below given blocks are accessed in this sequence:

W = 22, 11, 22, 44, 11, 33, 11, 22, 55, 27, 44
mccpp H H H H xcrv H

(M) cache miss - Must fetch
When block is not in the cache

(C) cold miss
when loc is available

(X) capacity miss
when cache is too small for
working set

(F) conflict miss
hen two or more blocks map to
same location

(H) cache hit
When block is in the cache so
data can me accessed

placement policies

P1 1. Unrestrictive

P2 2. Restrictive
block num % 16

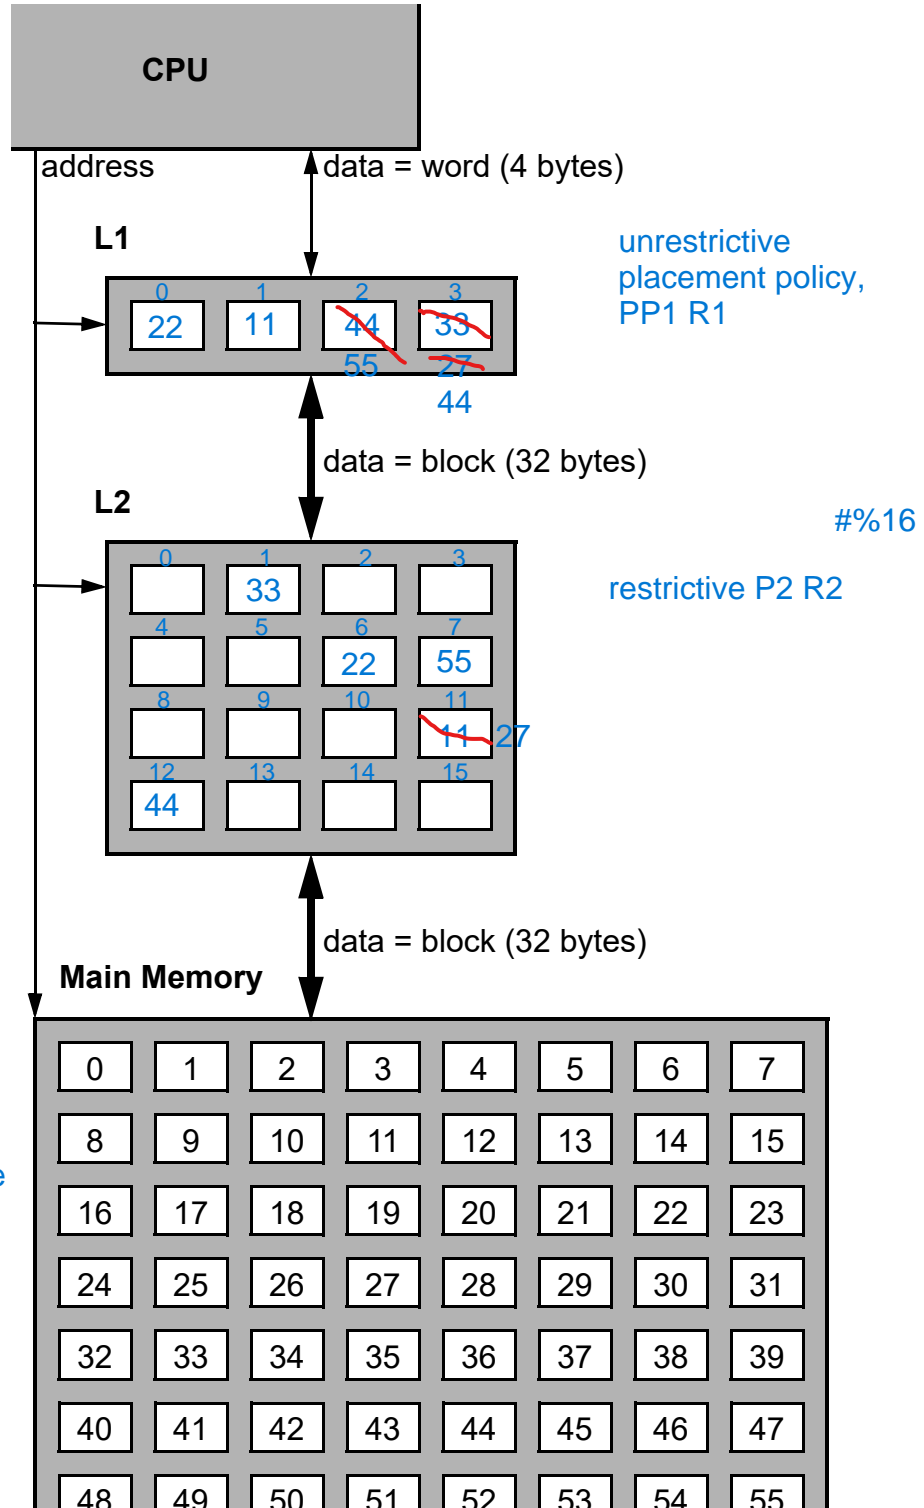
replacement policies

(R) 1. Choose any location

(R2) 2. No choice, because of restrictive
placement policy

(V) victim block
Cache block chosen for eviction

(W) working set
The set of blocks accessed during
some interval of time



Designing a Cache: Blocks

✱ *The bits of an address* are used to see if block containing address is in cache

IA-32 32 bits in address

How many bytes in an address space?


Let M be number of bytes in AS, IA-32 is 4GB

$$M = 2^m = 4 \text{ GB}$$

$$m = \log_2 M = \log_2(4\text{GB}) = 32 \text{ bits}$$

Thus m is number of bits in an address, IA-32 is 32 bits

32-bit Address Breakdown



How big is a block? # bytes

✱ *Cache blocks must be big enough* to capture spatial locality
but small enough to minimize latency

Let B be number of bytes per block, IA-32 is 32 bytes

$$B = 2^b = 32 \text{ bytes}$$

$$b = \log_2 B = \log_2(32) = 5 \text{ bits}$$

b bits: # of address bits needed to determine which byte in blocks

"word offset" identify which word in blocks 4 bytes/word 3 bits

"byte offset" identifies which byte in the word 2 bits

- What is the problem with using the most significant bits (left side) for the b bits?
Lose spatial locality, using least significant bits ensures contiguous block of addresses

How many 32-byte blocks of memory in a 32-bit address space?

$$\text{address space} / \text{block size} = 2^{32} / 2^5 = 2^{27} = 2^7 * 2^{20} = 128 \text{ Mb}$$

$$\text{recall: } K = 2^{10} \quad M = 2^{20} \quad G = 2^{30}$$

✱ *The remaining bits of an address* encode the block number.

Rethinking Addressing

* An address identifies **which byte in the VAS to access**.

* An address is **divided into parts to access memory into steps**


Memory Access in Caching System

step 1. Identify which **block in VAS**

step 2. Identify which **word in the block**


step 3. Identify which **byte in the word**

32-bit Address Breakdown



0xFFFFFFFF

IA-32 4GB Virtual Address Space



Designing a Cache: Sets & Tags

✱ *A cache must be searched* if unrestrictive placement policy is used

→ Problem? SLOW $O(N)$ where N is number of blocks in cache

Improvement? Limit or restrict location for each block

"set:" where block can be found

✱ *The block number bits of an address* are divided into two parts


1. maps block to specific set
2. a unique label or "tag" for the block

How many sets in the cache?

Let S be the number of sets in cache

$$S = 2^s = 1024 \text{ sets in cache}$$

$$s = \log_2 S = \log_2(1024) = 10$$



s bits: identify which set a block belongs to

➤ What is the problem with using the most significant bits (left side) for the s bits?

lose spatial locality


choose bits next to b -bits to ensure neighboring blocks go into neighboring sets

→ How many blocks map to each set for a 32-bit AS and a cache with 1024 sets? 8192 sets?

$$\text{address space} / \# \text{ sets} = 2^{(32 - 5)} / 2^{10} = 2^{17} = 128 \text{ k blocks mapping to same set}$$

Since different blocks map to the same set
how do we know which block is in a set?

use remaining bits as a unique tag "label"



t bits: bits of addr that identify block

✱ *When a block is copied into a cache its t bits are also stored as its tag*


Basic Cache Lines

What? A line is

- ◆ a location in a cache that stores one block of memory
- ◆ composed of storage bits for block and info needed by cache

- ✱ In our basic cache each cache set has only one line

Basic Cache Diagram



→ How do you know if a line in the cache is used or not?

use a status bit

v-bit: 1 = valid 0 = not valid

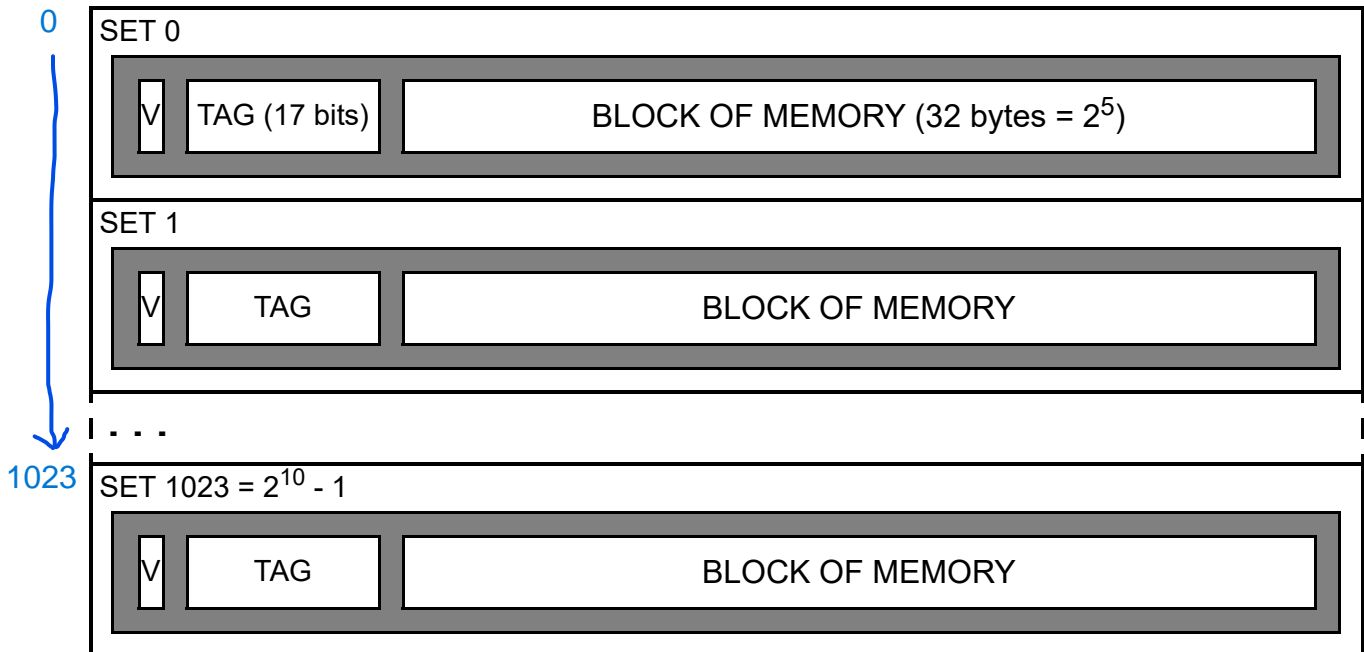
→ How big is a basic cache given S sets with blocks having B bytes?

cache size = $C = S \times B = \# \text{ sets times number of bytes per set}$

$$\# S \times \frac{\# \text{bytes}}{\text{Set}} = \text{_____ bytes}$$

Basic Cache Operation

Basic Cache Diagram



→ How big is this basic cache?

$$C = S \times B = 1024 \times 32 = 2^{10} \times 2^5 = 32 \text{ k}$$

How does a cache process a request for a word at a particular address?

1. Set Selection

extract S bits
use as index

2. Line Matching

extract t bits
compare t bits with stored tag

if no match or valid bit is 0, **cache MISS!**

must fetch block from next lower level

if match and valid bit is 1, **cache HIT**


for L1 cache, must now extract word from block

32-bit Address Breakdown



Basic Cache Practice

You are given the following 32-bit address breakdown used by a cache:



→ How big are the blocks?

6 b-bits $B = 2^b = 2^6 = 64$ bytes

→ How many sets?

13 s-bits $S = 2^s = 2^{13} = 8 \text{ k} = 8192$ sets


→ How big is this basic cache?

$C = S \times B = 2^{13} \times 2^6 = 2^{19} = 2^{10} \times 2^9 = 512 \text{ k}$

Assume the cache design above is given the following specific address: 0x07515E2B



→ Which set should be checked given the address above?



→ Which word in the block does the L1 cache access for the address?



➤ Which byte in the word does the address specify?

byte 3

Assume address above maps to a set with its line having the following V status and tag.

→ Does the address above produce a hit or miss?

V tag

1.) 1 0x0750 miss

2.) 0 0x0750 miss

3.) 1 0x00EA HIT

4.) 0 0x00EA miss

