

CS 354 - Machine Organization & Programming

Tuesday April 11, Thursday April 13, 2023

Exam Results expected by Friday April 14

Homework hw5DUE Monday 4/10**Homework hw6:** DUE on or before Monday Apr 17

Homework hw7: DUE on or before Monday Apr 24

Project p5: DUE on or before Friday April 21

Last Week

Instructions - SET Instructions - Jumps Encoding Targets Converting Loops	The Stack from a Programmer's Perspective The Stack and Stack Frames Instructions - Transferring Control Register Usage Conventions Function Call-Return Example
--	--

This Week

Function Call-Return Example (L20 p7) Recursion Stack Allocated Arrays in C Stack Allocated Arrays in Assembly Stack Allocated Multidimensional Arrays	Stack Allocated Structs ... rest was next week f22 Alignment Alignment Practice Unions
Next Week: Pointers in Assembly, Stack Smashing, and Exceptions B&O 3.10 Putting it Together: Understanding Pointers 3.12 Out-of-Bounds Memory References and Buffer Overflow 8.1 Exceptions 8.2 Processes 8.3 System Call Error Handling 8.4 Process Control through p719	

Recursion

Use a stack trace to determine the result of the call `fact(3)`:

```
int fact(int n) {
    int result;
    if (n <= 1) result = 1; //BASE CASE
    else      result = n * fact(n - 1); //REC CASE
    return result;
}
```

direct recursion when a function calls itself

recursive case when calls recursive function

base case stops recursion

"infinite" recursion similar to infinite loop, occurs when we don't reach a base case (eventually causes stack overflow, has to end)

%eax 4244446
%ebx 88888?

Assembly Trace

fact:

```
pushl %ebp
movl %esp, %ebp
pushl %ebx
subl $4, %esp
```

S.F. setup

```
movl 8(%ebp), %ebx - 1st arg
movl $1, %eax - 1
cmpl $1, %ebx ebx - 1, ZF = 0, SF = 0
jle .L1
```

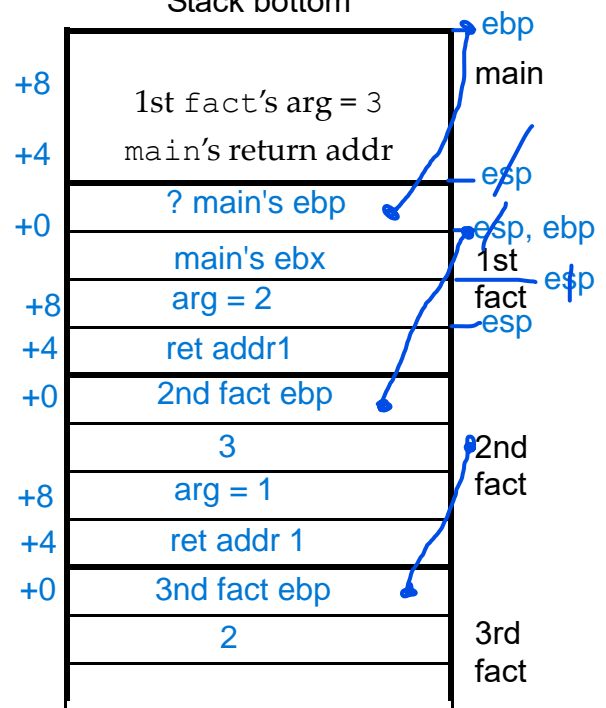
```
leal -1(%ebx), %eax
movl %eax, (%esp) - set arg
call fact push ret addr
jmp fact
```

ret addr 1 `imull %ebx, %eax` `eax <- eax * ebx`

```
.L1:
addl $4, %esp
popl %ebx
popl %ebp
ret popl %eip
```

MEM

Stack bottom



* "Infinite" recursion causes

* When tracing functions in assembly code

Stack Allocated Arrays in C

Recall Array Basics

$T \ A[N];$ where T is the element datatype of size L bytes and N is the number of elements



\swarrow $\text{sizeof}(T)$

1. contiguous region of stack mem $L * N$
2. identifies assoc. with start addr of array

* *The elements of A are accessed using address arithmetic*

Recall Array Indexing and Address Arithmetic

$$\&A[i] \equiv A + i \equiv x_A + L * i$$

\swarrow index
 \swarrow element size
 \swarrow start addr of A

→ For each array declarations below, what is L (element size), the address arithmetic for the i th element, and the total size of the array?


C code	L	address of i th element	total array size
* 1. <code>int I[11]</code>	4	$x_I + i * 4$	44
2. <code>char C[7]</code>			
3. <code>double D[11]</code>			
* 4. <code>short S[42]</code>	2	$x_S + 2 * i$	84
5. <code>char *C[13]</code>			
* 6. <code>int **I[11]</code>	4	$x_I + 4 * i$	44
7. <code>double *D[7]</code>			

Stack Allocated Arrays in Assembly

Arrays on the Stack

→ How is an array laid out on the stack? Option 1 or 2:

* *The first element (index 0) of an array is closest to top of stack*



Accessing 1D Arrays in Assembly

IA-32 Addr modes are designed to simplify array access

Assume array's start address in %edx and index is in %ecx

$\text{movl } (\%edx, \%ecx, 4), \%eax$
 $\begin{matrix} x_A & i & L \\ & & \text{(size of elt)} \end{matrix} \equiv m[x_A + 4 * i] \equiv *(A+i) \equiv A[i]$

→ Assume I is an `int` array, S is a `short int` array, for both the array's start address is in %edx, and the index i is in %ecx. Determine the element type and instruction for each:


C code	type	assembly instruction to move C code's value into %eax
* 1. I	<code>int*</code>	x_I <code>movl %edx, %eax</code>
* 2. I[0]	<code>int</code>	$m[x_I]$ <code>movl (%edx), %eax</code>
3. *I		
4. I[i]		$i \quad L$
* 5. &I[2]	<code>int*</code>	$x_I + 2 * 4$ <code>leal (%edx, %ecx, 4), %eax</code>
* 6. I+i-1	<code>int*</code>	$x_I + (i - 1)4$ <code>leal -4(%edx, %ecx, 4), %eax</code>
7. *(I+i-3)		
8. S[3]		
9. S+1		
10. &S[i]		
* 11. S[4*i+1]	<code>short</code>	$m[x_S + (4 * i + 1)2]$ <code>movswl 2(%edx, %ecx, 8), %eax</code>
12. S+i-5		

***valid scale factors include 1, 2, 4, and 8

Stack Allocated Multidimensional Arrays

Recall 2D Array Basics

$T \ A[R][C]$; where T is the element datatype of size L bytes,
 R is the number of rows and C is the number of columns



✱ Recall that 2D arrays are stored on the stack in row major order

```
int A[5][3];          typedef int row_t[3];    row_t is 10 array of 3 int
                      row_t A[5];             5 row_t array elements
```

Accessing 2D Arrays in Assembly

$\&A[i][j] \equiv$ Start addr + offset to row i + offset to col j
 $X_A + (L * C * i) + L * j$

Given array A as declared above, if x_A in $\%eax$, i in $\%ecx$, j in $\%edx$
then $A[i][j]$ in assembly is:

```
leal ( $\%ecx$ ,  $\%ecx$ , 2),  $\%ecx$      $i + 2i = 3i$ 
sall $2,  $\%edx$                    $4 * j$ 
addl  $\%eax$ ,  $\%edx$                 $x_A + 4j$ 
movl ( $\%edx$ ,  $\%ecx$ , 4),  $\%eax$      $x_A + 3i * 4 + 4j$ 
```

$$x_A + 12i + 4j$$

Compiler Optimizations

- ♦ If only accessing part of array then compiler makes a pointer to that part of array, and then compute offset from there
- ♦ If taking a fixed stride through the array the compiler uses $\text{stride} * L$ as offset

Stack Allocated Structures

Structures on the Stack

```

struct iCell {
    int x;
    int y;
    int c[3];
    int *v;
};
    
```

(hex)


#bytes	offset
4	0x0
4	0x4
12	0x8
4	0x14

→ How is a structure laid out on the stack? Option 1 or 2:

The compiler

◆ associates data member names with their offset from start of struct

◆ uses addr arith. with offsets to access data members



* The first data member of a structure is closest to top of stack

Accessing Structures in Assembly

Given:

```

struct iCell ic = //assume ic is initialized

void function(iCell *ip) {
    %esp
    
```

→ Assume `ic` is at the top of the stack, `%edx` stores `ip` and `%esi` stores `i`.

Determine for each the assembly instruction to move the C code's value into `%eax`:

C code	assembly
1. <code>ic.v</code>	<code>movl 20(%esp), %eax</code>
2. <code>ic.c[i]</code>	<code>movl 4(%esp, %esi, 4), %eax</code>
3. <code>ip->x</code>	<code>movl (%edx), %eax</code>
4. <code>ip->y</code>	<code>movl 4(%edx), %eax</code>
5. <code>&ip->c[i]</code>	<code>leal 8(%edx, %esi, 4), %eax</code>

* Assembly code to access a structure does not have data member names & types


Alignment

What? Most computer systems restrict addresses where primitive data can be stored

Why? Better memory performance

Example: Assume cpu reads 8 byte words
f is a misaligned float

Slower: requires two reads, extract bytes
recombine into single float value



Restrictions

IA-32 has NO alignment restriction


Linux: short addr must be mult of 2
int, float, pointer, double addr must be mult of 4

Windows: same as Linux except
double must be mult of 8

Implications padding might be inserted by compiler
into structs to keep data members aligned

Structure Example

```
struct s1 {  
    int i;  
    char c;  
    int j;  
};
```




✳ *The total size of a structure* is typically a multiple of size of longest data member

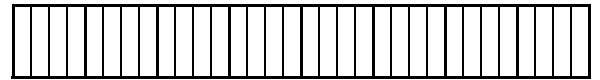
Alignment Practice

→ For each structure below, complete the memory layout and determine the total bytes allocated.


```
1) struct sA {  
    int i;  
    int j;  
    char c;  
};
```




```
2) struct sB {  
    char a;  
    char b;  
    char c;  
};
```




```
3) struct sC {  
    char c;  
    short s;  
    int i;  
    char d;  
};
```



```
4) struct sD {  
    short s;  
    int i;  
    char c;  
};
```



```
5) struct sE {  
    int i;  
    short s;  
    char c;  
};
```



✱ *The order that a structure's data members are listed* can affect mem. utilization because of alignment padding

Unions

What? A union is


- ♦ like a struct, except fields share same memory
bypassing C's type checking
- ♦ allocates only enough memory for the largest field

Why?

- ♦ allows data to be accessed as different types
- ♦ used to access hardware
- ♦ low-level polymorphism


How?

```
struct s {  
    char c;  
    int i[2];  
    double d;  
};
```



Total Size = 24 bytes


```
union u {  
    char c;  
    int i[2];  
    double d;  
};
```



Total Size = 8 bytes

Example

```
typedef union {  
    unsigned char cntrlrByte;  
    struct {  
        unsigned char playbutn : 1;  
        unsigned char pausebutn : 1;  
        unsigned char ctrlbutn : 1;  
        unsigned char fire1butn : 1;  
        unsigned char fire2butn : 1;  
        unsigned char direction : 3;  
    } bits;  
} CntrlrReg;
```



```
Cntrlr Reg r1; //1 byte  
readCntrlrReg(r1.CntrlrByte)  
if (r1.bits.fire1butn) { .... fire shot}
```