

# CS 354 - Machine Organization & Programming

## Tuesday April 25 , and Thursdat April 27th, 2023

**Homework hw7:** DUE on or before Monday Apr 24th

**Homework hw8:** DUE on Monday May 1st

**Homework hw9:** DUE on Wednesday May 3rd

**Project p6:** Due on last day of classes, May 5th. **Please complete p6 by Friday of this week as labs are very busy last week of classes.**

If you do plan on getting help during last week of classes, be sure to bring your own laptop in case there is no workstation available.

### Last Week

Pointers Function Pointers Buffer Overflow & Stack Smashing Flow of Execution Exceptional Events Kinds of Exceptions	Transferring Control via Exception Table Exceptions/System Calls in IA-32 & Linux Processes and Context User/Kernel Modes Context Switch Context Switch Example
---	--

### This Week

Meet Signals Three Phases of Signaling Processes IDs and Groups Sending Signals Receiving Signals	Issues with Multiple Signals Forward Declaration Multifile Coding Multifile Compilation Makefiles
<b>Next Week:</b> Linking and Symbols B&O 7.1 Compiler Drivers 7.2 Static Linking 7.3 Object Files 7.4 Relocatable Object Files 7.5 Symbols and Symbols Tables 7.6 Symbol Resolution 7.7 Relocation	

# Meet Signals

✱ *The Kernel uses signals to notify User processes of exceptional events.*

**What?** A signal is a small message sent to process via kernel

Linux: hardware (HW) has 30 standard signal types, each with unique non-negative ID#

\$kill -l see a list of signal names and numbers

signal(7) see man 7 signal

## Why?

- ◆ So kernel can notify user process about:
  1. low-level HW exceptions 0-31
  2. high-level events in kernel or user processes
- ◆ enable user processes to communicate with each other
- ◆ to implement high-level software form of exception handling

## Examples

1. divide by zero

exception #0 interrupts to kernel handler

- kernel signals user proc with SIGFPE #8 (floating point exception)

2. illegal memory reference

exception #13 interrupts to kernel handler

- kernel signals user proc with SIGSEGV #11

3. keyboard interrupt

- ctrl-c interrupts to kernel handler which signals SIG INT #2  
terminates foreground process (fg) by default

- ctrl-z interrupts to kernel handler which signals SIGTSTP #20 (suspend)  
suspends foreground process by default

# Three Phases of Signaling

## Sending

- ♦ when the kernel exception handler runs in response to an exceptional event or signal from user process
- ♦ is directed to a destination process


## Delivering

when the kernel records a sent signal for its destination process

pending signal is delivered but not received

- ♦ each process has a bit vector for recording pending signals

bit vectors kernel data structure where each bit has a distinct meaning



- ♦ bit k is set to 1 when that signal is pending (delivered)  
bit k is set to 0 when that signal is received

## Receiving

when the kernel causes destination process to react to pending signal

- ♦ Happens when kernel transfers control back to process
- ♦ multiple pending signals are in order from low to high signal numbers  
lower numbers have higher priority


blocking prevents a signal from being received

- ♦ enables process to control which signals it pays attention to
- ♦ each process has a second bit vector for blocking signals  
4 bytes      1 = block

# Process IDs and Groups

**What?** Each process

- ◆ is identified by a PID (a process ID number)
- ◆ and belongs to one group identified by PGID (process group ID number)



**Why?**

numbers pid are easier to manage than names  
pgid

**How?**

Recall: `ps` list running processes with pid & pgid

`jobs` list processes using simple job number

`getpid(2)`/`getpgrp(2)`

```
#include <unistd.h>
```

```
pid_t getpid(void) returns pid
```

```
pid_t getpgrp(void) returns pgid
```

# Sending Signals

**What?** A signal is sent by the kernel or a user process via the kernel [from command line](#)  
or in a program using system calls

## How? Linux Command

kill(1) [man 1 kill](#)      [sends signal from command line](#)  
[kill -l](#)                      [to specific pid](#)

kill -9 <pid>    [9 SIGKILL](#)              [2 SIGINT](#)              [20 SIGTSTP](#)

→ What happens if you kill your shell? [log out, same as exit](#)

## How? System Calls [with a program](#)

kill(2) [man 2 kill](#)      [send sig from calling to another process](#)

killpg(2) [send signal to all processes in process group](#)

```
#include <sys/types.h>
#include <signal.h>
int kill (pid_t pid, int sig)
```

[pid - process id of recipient & signal](#)  
[sig - signum](#)

[returns 0 on success](#)  
[-1 on error](#)

alarm(2) [man 2 alarm](#)  
[sets alarm that will deliver signal SIGALRM](#)

```
#include <unistd.h>
unsigned int alarm(unsigned int seconds)
```

[# sec remaining if prev alarm](#)  
[otherwise returns 0](#)

[until alarm is triggered](#)

# Receiving Signals

**What?** A signal is received by its destination process *by doing default action*  
*or executing code specified by sig handler*

## How? Default Actions

- ♦ Terminate the process *SIGINT #2 ctrl-c*
- ♦ Terminate the process and dump core *SIGSEGV #11*
- ♦ Stop the process *SIGSTP #20 ctrl-z*
- ♦ Continue the process if it's currently stopped *SIGCONT #18*
- ♦ Ignore the signal *SIGWINCH #28*

*(not the only signals  
with these default behaviors)*

## How? Signal Handler

### 1. *code a signal handler*

- ♦ *looks like a regular function but is called by kernel*
- ♦ *should not make unsafe system calls (like printf)*

### 2. *register the signal handler*

- ♦ *catch one or more signals*

~~signal(2)~~ **DON'T USE!**

sigaction(2) *POSIX for examining and changing a signal's action*

## Code Example

```
#include <signal.h>
#include ...
#include <string.h>

void handler_SIGALRM() { ... }

int main(...) {
    ...
    //Register our signal handler
    struct sigaction sa;
    memset(&sa, 0, sizeof(sa));
    sa.sa_handler = handler_SIGALRM;
    if (sigaction(SIGALRM, &sa, NULL) != 0) {
        printf("Error binding SIGALRM handler\n");
        exit(1);
    }
}
```

*code to handle alarm* (points to handler\_SIGALRM)

*function pointer* (points to sa.sa\_handler)

# Issues with Multiple Signals

**What?** Multiple signals of the same type as well as those of different types can be sent during same period that other signals are sent and even while sig. handler is running

## Some Issues

→ Can a signal handler be interrupted by other signals?  
yes, but ... linux signals of same type don't interrupt  
they become pending

✱ *Block any signals* that you don't want to interrupt your sig handler  
sigemptyset(&sa.sa\_mask); // blocks all  
sigfullset(&sa.sa\_mask); // enable all (unblock all)  
sigaddset, sigdelset, sigismember(&sa.sa\_mask, signum);

→ Can a system call be interrupted by a signal? yes, for...  
slow system calls which potentially take a long time  
such system calls return immediately EINTR  
sa.sa\_flags = SA\_RESTART; //attempt to restart

NOTE: sleep() can not be restarted

→ Does the system queue multiple standard signals of the same type for a process? **NO**  
the bit vector indicates pending or NOT pending, THERE IS NO COUNT

✱ *Your signal handler shouldn't assume* that a signal was sent only one time (except in p6)

## Real-time Signals

Linux has 33 additional application defined signals

- ◆ They can include integer or pointer in their message
- ◆ Multiple signals of same type are queued in the order delivered
- ◆ Multiple signals of different types are received from low to high sig number

//all necessary content for p6 now covered, also in textbook

# Forward Declaration

**What?** Forward declaration tells compiler about certain attributes of identifier before it's fully defined

✱ *Recall, C requires that an identifier* be declared before used

**Why?** gcc is a

- ♦ one pass compiler (gcc) can ensure identifier exists and is correctly used
- ♦ large programs can be divided into separate functional units and compile separately
- ♦ mutual recursion is possible. A calls B  
B calls A

## Declaration vs. Definition

declaring tells compiler about

variables: name and type    `extern int g;`

functions: return type, name, parameter types & order

defining provides full details

variables: where in memory it's allocated

functions: function body is defined

✱ *Variable declarations* usually both declare and define

```
void f(){  
    int i = 11;    // declare, define, & initialize  
    static int j;  // declaration and definition (not initialized)
```

✱ *A variable is proceeded with* `extern` is NOT DEFINED

`extern char * Title; //declaration only`



# Multifile Coding

**What?** Multifile coding divides a program into functional units  
each unit has its own header and source file

## Header File (filename.h) - "public" interface

contains things intended to share  
mainly function declarations  
but also can have types, constants, macros

p3Heap.h

recall ~~heapAlloc.h~~ from project p3:

```
#ifndef __heapAlloc_h__
#define __heapAlloc_h__

int  initHeap(int sizeOfRegion);
void* allocHeap(int size);
int  freeHeap(void *ptr);
void dumpMem();

#endif // __heapAlloc_h__
```

function declaration

\* *An identifier* can be defined only once in global scope

#include guard: prevents multiple inclusions of the same header file

## Source File (filename.c) - "private" implementation

must include definitions of things in header file

p3Heap.c

recall ~~heapAlloc.c~~ from project p3:

```
#include <unistd.h>
...
#include "heapAlloc.h"

typedef struct blockHeader {
    int size_status;
} blockHeader;

blockHeader *heapStart = NULL;

void* allocHeap(int size) { . . . }
int  freeHeap(void *ptr) { . . . }
int  initHeap(int sizeOfRegion) { . . . }
void dumpMem() { . . . }
```

"private"

defns

# Multifile Compilation

**gcc Compiler Driver** directs tools needed to build executable



**Object Files** binary code & data segment

relocatable object file (ROF) can be combined by the linker with other ROFs and SOFs to produce an EOF

executable object file (EOF) can be loaded into mem and "run"

shared object file (SOF) can be loaded into memory (DLL) or statically linked into EOF

## Compiling All at Once

```
gcc align.c heapAlloc.c -o align -m32 -Wall -std=gnu99
```

src → EOF

## Compiling Separately

```
gcc -c align.c           compile to align.o (ROF)
gcc -c heapAlloc.c       compile o heapAlloc.o (ROF)
gcc align.o heapAlloc.o -o align link to align (EOF)
```

\* *Compiling separately is* more efficient and easier to manage

## What? Makefiles are

♦

♦

## Why?

♦

♦

## Rules

## Example

```
#simplified p3 Makefile
rule 1 align: align.o heapAlloc.o
    gcc align.o heapAlloc.o -o align
rule 2 align.o: align.c
    gcc -c align.c
rule 3 heapAlloc.o: heapAlloc.c heapAlloc.h
    gcc -c heapAlloc.c
rule 4 clean:
    rm *.o
    rm align
```

## Using

```
$ls
align.c Makefile heapAlloc.c heapAlloc.h
$make
gcc -c align.c
gcc -c heapAlloc.c
gcc align.o heapAlloc.o -o align
$ls
align align.c align.o Makefile heapAlloc.c heapAlloc.h heapAlloc.o
$rm heapAlloc.o
rm: remove regular file 'heapAlloc.o'? y
$make
gcc -c heapAlloc.c
gcc align.o heapAlloc.o -o align
$make heapAlloc.o
make: 'heapAlloc.o' is up to date.
$make clean
rm *.o
rm align
$ls
align.c Makefile heapAlloc.c heapAlloc.h
```